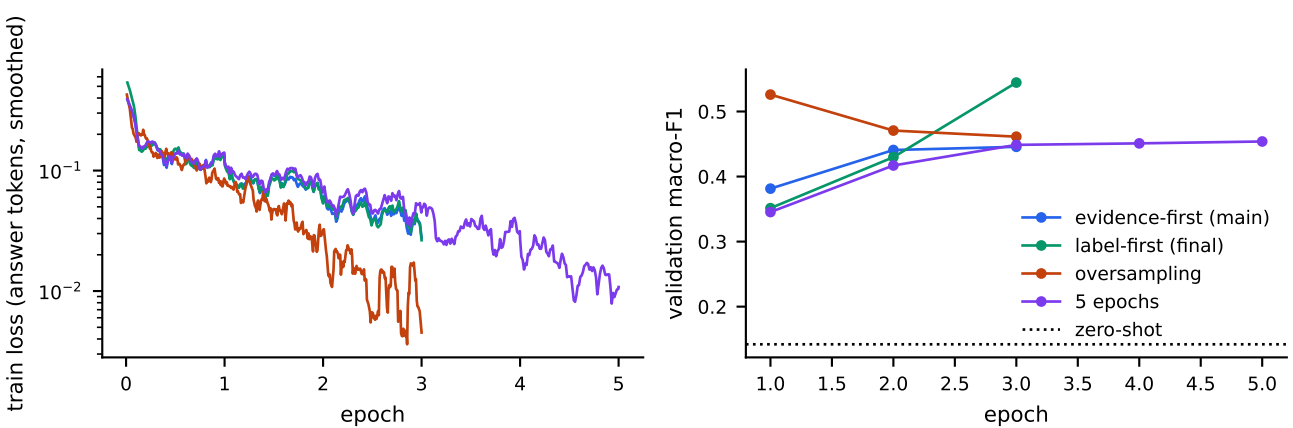
train loss (answer tokens, smoothed)
10^{-1}
10^{-2}
0
1
2
3
4
5
epoch
validation macro-F1
0.5
0.4
0.3
0.2
1.0
1.5
2.0
2.5
3.0
3.5
4.0
4.5
5.0
epoch
evidence-first (main)
label-first (final)
oversampling
5 epochs
zero-shot